

Chapter 8 examples

Brooke Anderson

4/7/2020

Load some general packages you'll need:

```
library(tidyverse)
```

Pasilla dataset

Install the **pasilla** package and find the file address of the example data file you'll be using. The **system.file** function looks for a file that was included with a package. If you want to share data using something other than an R binary data file, this is how the data is included in the package. The output (**fn**) will be the filepath to the file "pasilla_gene_counts.tsv", which you downloaded as part of the **pasilla** package.

```
# Uncomment and run the following line if you need the `pasilla` package
# BiocManager::install(pkgs = "pasilla")
fn <- system.file("extdata", "pasilla_gene_counts.tsv",
                  package = "pasilla", mustWork = TRUE)
fn
```

```
## [1] "/Library/Frameworks/R.framework/Versions/3.6/Resources/library/pasilla/extdata/pasilla_gene_cou"
```

Now that you know where the file is, you need to read it in:

```
counts <- fn %>%
  read_tsv()
counts
```

```
## # A tibble: 14,599 x 8
##   gene_id untreated1 untreated2 untreated3 untreated4 treated1 treated2
##   <chr>          <dbl>      <dbl>      <dbl>      <dbl>      <dbl>      <dbl>
## 1 FBgn00~         0         0         0         0         0         0
## 2 FBgn00~        92        161        76        70        140        88
## 3 FBgn00~         5         1         0         0         4         0
## 4 FBgn00~         0         2         1         2         1         0
## 5 FBgn00~       4664       8714       3564       3150       6205       3072
## 6 FBgn00~        583        761        245        310        722        299
## 7 FBgn00~         0         1         0         0         0         0
## 8 FBgn00~        10        11         3         3         10         7
## 9 FBgn00~         0         1         0         0         0         1
## 10 FBgn00~      1446      1713        615        672      1698        696
## # ... with 14,589 more rows, and 1 more variable: treated3 <dbl>
```

Again, this data is “transposed” compared to what we’ll want for a lot of statistical modeling functions. Each column is an observation, while each row is the measures for a specific gene (which one is given in the first column, **gene_id**).

There is a vignette for the **pasilla** package at: https://bioconductor.org/packages/release/data/experiment/vignettes/pasilla/inst/doc/create_objects.html The paper that the data originally came from is available

here.

You can explore the data a bit:

```
# See how many rows and columns the data has
counts %>%
  dim()
```

```
## [1] 14599      8
```

```
# See a summary of each column
counts %>%
  summary()
```

```
##      gene_id      untreated1      untreated2      untreated3
## Length:14599    Min.   :    0.0    Min.   :    0    Min.   :    0.0
## Class :character 1st Qu.:    0.0    1st Qu.:    1    1st Qu.:    0.0
## Mode  :character Median :   27.0    Median :   46    Median :   18.0
##              Mean  :  957.1    Mean  : 1501    Mean  :   572.5
##              3rd Qu.:  637.0    3rd Qu.:  990    3rd Qu.:  351.0
##              Max.   :232141.0    Max.   :360330    Max.   :131242.0
##      untreated4      treated1      treated2      treated3
## Min.   :    0.0    Min.   :    0.0    Min.   :    0.0    Min.   :    0.0
## 1st Qu.:    0.0    1st Qu.:    1.0    1st Qu.:    0.0    1st Qu.:    0.0
## Median :   19.0    Median :   47.0    Median :   20.0    Median :   22.0
## Mean   :   674.1    Mean   : 1278.9    Mean   :   655.6    Mean   :   708.5
## 3rd Qu.:  413.5    3rd Qu.:  902.5    3rd Qu.:  416.0    3rd Qu.:  456.0
## Max.   :167116.0    Max.   :253500.0    Max.   :146390.0    Max.   :164148.0
```

```
# See a random sample of 5 rows
counts %>%
  sample_n(size = 5)
```

```
## # A tibble: 5 x 8
##   gene_id untreated1 untreated2 untreated3 untreated4 treated1 treated2 treated3
##   <chr>      <dbl>      <dbl>      <dbl>      <dbl>      <dbl>      <dbl>      <dbl>
## 1 FBgn00~      0          0          0          0          0          0          0
## 2 FBgn00~      3          0          1          1          0          2          0
## 3 FBgn00~      1          6          1          1          6          2          3
## 4 FBgn00~      0          0          0          0          0          0          0
## 5 FBgn00~      0          2          0          1          0          0          0
```

```
# Check out the unique letters in the gene names
# (i.e., once you take out all the digits at the end)
```

```
counts %>%
  pull(gene_id) %>% # Extract just the `gene_id` column as a vector
  str_remove("[0-9].+") %>% # Remove the first digit and anything after
                             # using regular expressions
  unique() # Look at just the unique values
```

```
## [1] "FBgn"
```

It looks like every gene id starts with “FBgn” and then some 7-digit numeric ID.

Next, it might be helpful to sample a few of the genes and then visualize how the values recorded for each gene vary across the samples, showing treated / untreated with color:

```
counts %>%
  sample_n(size = 20) %>%
```

```


pivot_longer(cols = -gene_id, names_to = "sample_id", values_to = "value") %>%
mutate(treatment = str_remove(sample_id, "[0-9]")) %>% # Extract untreated / treated
# from sample IDs by getting
# rid of the number in each

ggplot(aes(x = treatment, y = value)) +
labs(x = "", y = "") +
geom_point(alpha = 0.5, size = 1) + # Use some transparency (alpha) to see if points
# are plotted on top of each other.

facet_wrap(~ gene_id, scales = "free_y") + # Since scales are very different across
# genes, set scales as free on y-axis

# Rotate the x-axis labels a bit (otherwise, they overlap)
theme(axis.text.x = element_text(angle = 45, hjust = 1))

```



Check out how much the range in the measured levels for genes varies across genes:

```

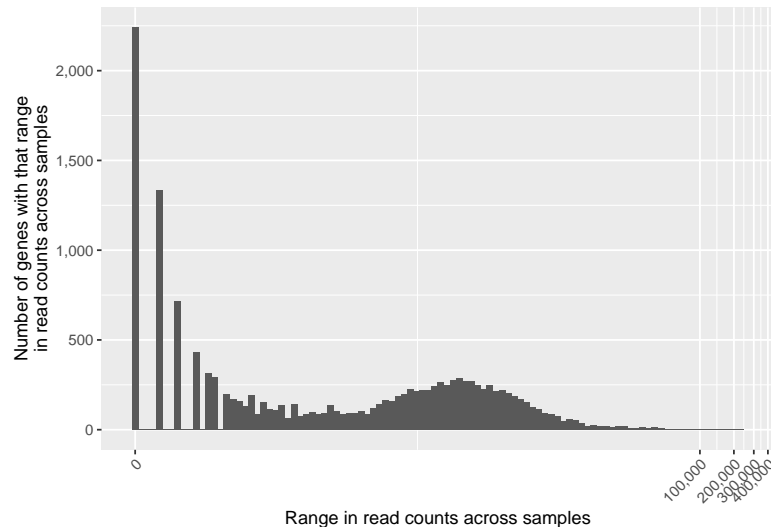
library("scales")
counts %>%
# Rearrange the data shape to making grouping and summarizing easier
pivot_longer(-gene_id, names_to = "sample_id", values_to = "value") %>%
# Group by gene and then calculate the minimum and maximum levels of each gene
# and the difference between those (range)
group_by(gene_id) %>%
summarize(min_value = min(value),
          max_value = max(value),
          range = max_value - min_value) %>%
# Plot a histogram of these ranges
ggplot(aes(x = range)) +
geom_histogram(bins = 100) +

```

```

labs(x = "Range in read counts across samples",
     y = "Number of genes with that range\nin read counts across samples") +
# Use a transformation on the x-axis---otherwise, with a few outliers, it's
# hard to see variation among the low values. You can't just use "log" in this
# case because you have a lot of zero values.
scale_x_continuous(trans = "pseudo_log", labels = comma) +
scale_y_continuous(labels = comma) +
theme(axis.text.x = element_text(angle = 45, hjust = 1))

```




Check typical count levels (for genes with highest levels) by sample.

```

library(purrr)
library(viridis)
counts %>%
  pivot_longer(-gene_id, names_to = "sample_id", values_to = "value") %>%
  # Get top 40 genes in terms of mean expression
  group_by(gene_id) %>%
  nest() %>%
  mutate(mean = map_dbl(data, ~ mean(.x$value))) %>%
  ungroup() %>%
  top_n(n = 40) %>%
  select(-mean) %>%
  unnest(data) %>%
  ggplot(aes(x = sample_id, y = gene_id, fill = value)) +
  geom_tile() +
  scale_fill_viridis()

```



Read in the meta-data for these data. The meta-data are available in a different file in the `padilla` package (“`pasilla_sample_annotation.csv`”).

```
annotationFile <- system.file("extdata", "pasilla_sample_annotation.csv",
                              package = "pasilla", mustWork = TRUE)
pasillaSampleAnno <- annotationFile %>%
  read_csv()
pasillaSampleAnno
```

```
## # A tibble: 7 x 6
##   file      condition type      `number of lane~` `total number of r~` `exon counts`
##   <chr>      <chr>      <chr>          <dbl> <chr>              <dbl>
## 1 treated~ treated   single~         5 35158667           15679615
## 2 treated~ treated   paired~         2 12242535 (x2)         15620018
## 3 treated~ treated   paired~         2 12443664 (x2)         12733865
## 4 untreat~ untreated single~         2 17812866           14924838
## 5 untreat~ untreated single~         6 34284521           20764558
## 6 untreat~ untreated paired~         2 10542625 (x2)         10283129
## 7 untreat~ untreated paired~         2 12214974 (x2)         11653031
```

This file gives some general information about each sample. This includes the file name where the sample’s data was saved (probably from the sampling equipment), the treatment or condition (treated / untreated), the type of read (single-read or paired-end), the number of lanes, total number of reads, and exon counts.

It can be hard to work with column names with spaces in the middle, so I recommend renaming the columns (replacing every space with an underscore will work):

```
# Current column names (they have spaces)
pasillaSampleAnno %>%
```

```
colnames()

## [1] "file"          "condition"      "type"
## [4] "number of lanes" "total number of reads" "exon counts"

# Rename columns to get rid of spaces
pasillaSampleAnno <- pasillaSampleAnno %>%
  # \\s stands for a space in R regular expressions
  rename_all(.funs = str_replace_all, "\\s", "_")

# New column names (no more spaces)
pasillaSampleAnno %>%
  colnames()

## [1] "file"          "condition"      "type"
## [4] "number_of_lanes" "total_number_of_reads" "exon_counts"
```

We can see some summaries of the data by condition (treated / untreated) and type (single-read / pair-ended):

```
pasillaSampleAnno %>%
  # Group the data and use `count` to count the number of rows with each combo
  group_by(condition, type) %>%
  count() %>%
  ungroup() %>%
  # Reshape the data to make a prettier table
  pivot_wider(names_from = type, values_from = n) %>%
  # Use `kable` from the `knitr` package to output as a pretty table
  knitr::kable()
```

condition	paired-end	single-read
treated	2	1
untreated	2	2

Create a `DESeqDataSet` object with both the measurements (from `counts`) and the meta-data (from `pasillaSampleAnno`). This *data container* is a special type of object class that will help keep everything in order as you move through the analysis (and helps with managing the size of the data).

```
library("DESeq2")

# Need to remember what the column names of `counts` are (so we can
# use the same ones for the meta-data)
colnames(counts)

## [1] "gene_id"      "untreated1" "untreated2" "untreated3" "untreated4"
## [6] "treated1"     "treated2"   "treated3"

# Clean up the meta-data a bit before we add it to this special object class
pasillaSampleAnno <- pasillaSampleAnno %>%
  rename(file = "sample_id") %>%
  # Use regular expressions to clean up sample_id name (so it matches
  # the column names in `counts`)
  mutate(sample_id = str_remove(sample_id, "fb")) %>%
  # Evidently, the "-"s in the `type` column may cause problems in DESeq,
  # so shorten those labels
  mutate(type = str_remove(type, "-.+")) %>%
  # Convert the `condition` and `type` columns to factors and
```

```

# set the level order by hand
mutate(condition = condition %>%
  as_factor() %>%
  fct_relevel("untreated", "treated"),
  type = type %>%
  as_factor() %>%
  fct_relevel("single", "paired")) %>%
# Finally, you need to make sure the rows of this dataframe
# are in the same order as the columns in the read count data matrix.
# One way to do this is by making the column names of the count data
# matrix (without the gene_id column) into a dataframe, with numbers
# giving the order, join that with this data, then reorder by that
# order number
full_join(counts %>%
  # Get the column names from `counts`
  colnames() %>%
  # Get rid of the first column name (which is the `gene_id` column)
  `[(-1) %>%
  # `enframe` will convert to a dataframe (from a vector)
  enframe(name = "order"),
  by = c("sample_id" = "value")) %>%
# Rearrange by order, then get rid of that column
arrange(order) %>%
select(-order)

# Here's what the meta-data looks like now:
pasillaSampleAnno

## # A tibble: 7 x 6
##   sample_id condition type    number_of_lanes total_number_of_reads exon_counts
##   <chr>         <fct>    <fct>          <dbl> <chr>                <dbl>
## 1 untreated1 untreated single           2 17812866             14924838
## 2 untreated2 untreated single           6 34284521             20764558
## 3 untreated3 untreated paired           2 10542625 (x2)        10283129
## 4 untreated4 untreated paired           2 12214974 (x2)        11653031
## 5 treated1    treated  single           5 35158667             15679615
## 6 treated2    treated  paired           2 12242535 (x2)        15620018
## 7 treated3    treated  paired           2 12443664 (x2)        12733865

# Put all the required components into a DESeqDataSet object so you can
# run DESeq on it
pasilla <- counts %>%
  # For the read data, it's all in counts, but we need to take off the
  # first column (with the gene IDs) and then convert to a matrix to
  # input it into this object class
  column_to_rownames("gene_id") %>%
  as.matrix() %>%
  # Put everything into a DESeqDataSet object
  DESeqDataSetFromMatrix(colData = pasillaSampleAnno,
    design = ~ condition)

class(pasilla)

## [1] "DESeqDataSet"
## attr(,"package")

```

```
## [1] "DESeq2"
```

DESeq

Apply DESeq algorithm to the data:

```
pasilla <- pasilla %>%  
  DESeq()
```

```
## estimating size factors
```

```
## estimating dispersions
```

```
## gene-wise dispersion estimates
```

```
## mean-dispersion relationship
```

```
## final dispersion estimates
```

```
## fitting model and testing
```

You can use *extractor functions* like `results` to pull out specific data from the resulting object. For example, you can pull out gene-specific estimates of differences between treated and untreated samples with:

```
pasilla %>%  
  results()
```

```
## log2 fold change (MLE): condition treated vs untreated
```

```
## Wald test p-value: condition treated vs untreated
```

```
## DataFrame with 14599 rows and 6 columns
```

```
##           baseMean      log2FoldChange      lfcSE  
##           <numeric>           <numeric>    <numeric>  
## FBgn0000003 0.171568715207063    1.02601368333522 3.80551160374507  
## FBgn0000008 95.1440789963134 0.00215175720349084 0.223883696572144  
## FBgn0000014 1.05657219346166 -0.496735176118498 2.16026588878143  
## FBgn0000015 0.846723274987709 -1.88276477012506 2.10643312162068  
## FBgn0000017 4352.5928987935 -0.240025038806395 0.126024438560778  
## ...           ...           ...           ...  
## FBgn0261571 0.087343676946538    0.90026855885989 3.81017301615324  
## FBgn0261572 6.19713652050888 -0.959130034959993 0.777016744083823  
## FBgn0261573 2240.98398636611 0.0126159820947047 0.112700631633334  
## FBgn0261574 4857.74267170989 0.0152573285663809 0.19314843580021  
## FBgn0261575 10.6835537573787 0.163562434452904 0.938909701933571  
##           stat           pvalue      padj  
##           <numeric>           <numeric>    <numeric>  
## FBgn0000003 0.269612548895004 0.787458345134478      NA  
## FBgn0000008 0.00961104911360736 0.99233161035707 0.996928202137134  
## FBgn0000014 -0.229941683890912 0.818137084970592      NA  
## FBgn0000015 -0.893816542666432 0.371420056652208      NA  
## FBgn0000017 -1.90459121696969 0.0568332291795643 0.282363564717396  
## ...           ...           ...           ...  
## FBgn0261571 0.236280230594043 0.813215226384258      NA  
## FBgn0261572 -1.23437498903695 0.217063204412214      NA  
## FBgn0261573 0.111942425804233 0.910869057061802 0.982036790151162  
## FBgn0261574 0.0789927627587049 0.937038379558298 0.988141534972512  
## FBgn0261575 0.174204648344848 0.861704632545053 0.967913643037464
```

If you want to work more easily with this output, you can convert it to a tibble with `as_tibble`:


```

pasilla %>%
  results() %>%
  as.data.frame() %>%
  rownames_to_column(var = "gene_id") %>%
  as_tibble()

```

```

## # A tibble: 14,599 x 7
##   gene_id      baseMean log2FoldChange lfcSE      stat pvalue   padj
##   <chr>         <dbl>          <dbl> <dbl>    <dbl> <dbl>   <dbl>
## 1 FBgn0000003    0.172            1.03   3.81    0.270  0.787   NA
## 2 FBgn0000008   95.1             0.00215 0.224  0.00961 0.992  0.997
## 3 FBgn0000014    1.06            -0.497  2.16   -0.230  0.818   NA
## 4 FBgn0000015    0.847            -1.88   2.11   -0.894  0.371   NA
## 5 FBgn0000017  4353.            -0.240  0.126  -1.90   0.0568  0.282
## 6 FBgn0000018   419.            -0.105  0.148  -0.707  0.480   0.824
## 7 FBgn0000022    0.0797          -0.720  3.81   -0.189  0.850   NA
## 8 FBgn0000024    6.41             0.211  0.690  0.305  0.760   NA
## 9 FBgn0000028    0.439             1.41   2.78   0.509  0.611   NA
## 10 FBgn0000032   990.            -0.0919 0.148  -0.622  0.534   0.850
## # ... with 14,589 more rows

```

Check out the top few genes in terms of adjusted p-values for the Wald test comparing the untreated versus treated samples. Since smaller p-values are more interesting, pick out the smallest:

```

pasilla %>%
  results() %>%
  as.data.frame() %>%
  rownames_to_column(var = "gene_id") %>%
  filter(!is.na(padj)) %>%
  as_tibble() %>%
  arrange(padj) %>%
  dplyr::slice(1:5)

```

```

## # A tibble: 5 x 7
##   gene_id      baseMean log2FoldChange lfcSE      stat      pvalue      padj
##   <chr>         <dbl>          <dbl> <dbl>    <dbl>    <dbl>    <dbl>
## 1 FBgn0039155    731.            -4.62  0.169  -27.4  4.89e-165  4.07e-161
## 2 FBgn0025111   1501.             2.90  0.127   22.8  1.53e-115  6.38e-112
## 3 FBgn0029167   3706.            -2.20  0.0970 -22.7  1.33e-113  3.69e-110
## 4 FBgn0003360   4343.            -3.18  0.144  -22.2  9.56e-109  1.99e-105
## 5 FBgn0035085    638.            -2.56  0.137  -18.6  1.29e- 77  2.14e- 74


```

Histogram of p-values:

```

pasilla %>%
  results() %>%
  as_tibble() %>%
  ggplot(aes(x = pvalue)) +
  geom_histogram(binwidth = 0.01, fill = "Royalblue", boundary = 0)

```



MA plot (but using ggplot instead of plotMA from DESeq2 package—the x-axis transform might not be identical...):

```
pasilla %>%
  results() %>%
  as_tibble() %>%
  mutate(low_padj = padj < 0.1) %>%
  mutate(log2FoldChange = ifelse(log2FoldChange > 2, 2, log2FoldChange),
         log2FoldChange = ifelse(log2FoldChange < -2, -2, log2FoldChange)) %>%
  ggplot(aes(x = baseMean, y = log2FoldChange, color = low_padj)) +
  geom_point(size = 0.5) +
  ylim(c(-2, 2)) +
  scale_color_manual(values = c("black", "red")) +
  scale_x_continuous(trans = "pseudo_log") +
  geom_hline(yintercept = 0, color = "red", size = 2, alpha = 0.4)
```

